

## TABLE OF CONTENTS

TITLE .....	2
INTRODUCTION .....	2
DATASETS .....	2
EXPLANATION AND PREPARATION OF DATASETS .....	3
<i>Exploratory Data Analysis</i> .....	5
<i>Feature Improvement</i> .....	8
IMPLEMENTATION IN PYTHON .....	9
<i>Brief description of the algorithms used</i> .....	9
K-Nearest Neighbors (KNN) algorithm .....	9
Artificial Neural Networks .....	9
<i>Applications of the Algorithms</i> .....	9
K-Nearest Neighbors Classification .....	9
Application of the KNN algorithm to the dataset & Model design .....	9
Experimental procedure .....	10
Result Visualization: .....	13
Relevant Literature .....	13
Artificial Neural Networks (ANNs) classification .....	14
Application of the ANN algorithm to the dataset .....	14
Explanation of the experimental procedure .....	14
Model Parameters Used .....	14
Evaluating the Neural Network.....	15
Relevant Literature .....	16
IMPLEMENTATION IN AZURE MACHINE LEARNING STUDIO .....	17
<i>Pre-processing</i> .....	17
<i>Pipeline Population</i> .....	19
<i>Results</i> .....	20
RESULTS ANALYSIS AND DISCUSSION .....	23
<i>Performance metric used</i> .....	23
<i>Presentation of results</i> .....	23
KNN Classification .....	23
Artificial Neural Network .....	24
Azure Machine Learning .....	25
<i>Comparison and discussion of results</i> .....	26
<i>Ethical, legal, and professional considerations</i> .....	27
CONCLUSION .....	27

## **Title**

The prediction and prevention of phishing attacks by the classification of webpages into fraudulent (phishing) and legitimate categories using classification algorithms.

## **Introduction**

Computers are the driving force responsible for the advancement of technology in the world today, and with their constantly increasing use comes the risk of cybercrimes. Ranking high among these cybercrimes is the scourge that is Phishing (Financial and cybercrimes top global police concerns, says new Interpol report, 2022)

Phishing is the attempt to manipulate a user into performing an action that leads to unwanted consequences such as downloading malicious software (malware) or redirection to a questionable website. (Phishing attacks: Defending your organisation, 2018)

While performing everyday online activities, different website links/webpages are encountered, and the ability to distinguish between phishing and authentic websites cannot be overemphasized.

Machine learning classification systems can help make the process of correctly identifying these websites automatic, and this can be implemented in the design of softwares that automatically predicts and prevents this cyber attack.

This report seeks to evaluate and compare the performance of models, fabricated using 2 classification algorithms (K-Nearest Neighbors and Neural Networks) on the dataset provided in the following section. These models are designed to classify websites into either phishing or legitimate categories.

## **Datasets**

The dataset was published on the 24th of March, 2018 by Choon Lin Tan, and is titled, "Phishing Dataset for Machine Learning: Feature Evaluation", and was downloaded from <https://data.mendeley.com/datasets/h3cgnj8hft/1>

The data has 10,000 observations, which consist of 50 attributes taken from 5000 authentic websites (from Alexa and Common Crawl) and 5000 fraudulent websites (from PhishTank and OpenPhish) between January & May 2015 and May & June 2017.

The contributor, while obtaining the data, utilised the Selenium WebDriver, which is a browser automation framework, which provides improved webpage feature extraction techniques that delivers more accurate and reliable results than a parsing approach based on regular expressions. (Tan, 2018).

The dataset was downloaded in .arff format and was converted to .csv through the use of an open-source Github online converter.

In the class column (CLASS\_LABEL), the phishing websites are represented by 1, the legitimate websites are represented by 0

## Explanation and preparation of datasets

The dataset was downloaded and read into a variable called 'dataset', using Pandas,

The dataset contains 50 attributes. The first attribute represents the Id, which will later be removed, the following 48 attributes are various features of the website in each observation, and the last attribute represents the class label, which is eventually renamed 'Result'. See below

```
In [120]: column_list = dataset.columns.values.tolist()
i_v = pd.DataFrame({'Independent Variables':column_list[:-1]})
i_v.index+=1
```

Independent Variables	
1	id
2	NumDots
3	SubdomainLevel
4	PathLevel
5	UrlLength
6	NumDash
7	NumDashInHostname
8	AtSymbol
9	TildeSymbol
10	NumUnderscore
11	NumPercent
12	NumQueryComponents
13	NumAmpersand
14	NumHash
15	NumNumericChars
16	NoHttps
17	RandomString
18	IpAddress
19	DomainInSubdomains
20	DomainInPaths
21	HttpsInHostname
22	HostnameLength
23	PathLength
24	QueryLength
25	DoubleSlashInPath
26	NumSensitiveWords
27	EmbeddedBrandName
28	PotExtHyperlinks
29	PotExtResourceUris
30	ExtFavicon
31	InsecureForms
32	RelativeFormAction
33	ExtFormAction
34	AbnormalFormAction
35	PotNullSelfRedirectHyperlinks
36	FrequentDomainNameMismatch
37	FakeLinkInStatusBar
38	RightClickDisabled
39	PopUpWindow
40	SubmitInfoToEmail
41	IframeOrFrame
42	MissingTitle
43	ImagesOnlyInForm
44	SubdomainLevelRT
45	UrlLengthRT
46	PotExtResourceUrisRT
47	AbnormalExtFormActionR
48	ExtMetaScriptLinkRT
49	PotExtNullSelfRedirectHyperlinksRT

Dependent variable:

```
In [50]: column_list = dataset.columns.values.tolist()
v = pd.DataFrame({'Dependent Variable':column_list[-1:]})
v.index+=1
v
```

```
Out[50]:
```

Dependent Variable	
1	CLASS_LABEL

Several preprocessing steps were followed.

The shape of the dataset was checked to see the number of records and attributes, then the '.head' function was used for a glance. See below

```
In [483]: dataset = pd.read_csv("C:\\Users\\C\\Downloads\\Phishing_Legitimate_full.csv")

In [484]: dataset.shape #to check the number of records and attributes

Out[484]: (10000, 50)

In [485]: dataset.head()

Out[485]:
```

	id	NumDots	SubdomainLevel	PathLevel	UrlLength	NumDash	NumDashInHostname	AtSymbol	TildeSymbol	NumUnderscore	...	IframeOrFrame	MissingT
0	1	3	1	5	72	0	0	0	0	0	...	0	
1	2	3	1	3	144	0	0	0	0	2	...	0	
2	3	3	1	2	58	0	0	0	0	0	...	0	
3	4	3	1	6	79	1	0	0	0	0	...	0	
4	5	3	0	4	46	0	0	0	0	0	...	1	

5 rows x 50 columns

Next, the '.info' function was used as a quick means to check for missing values and data type, and the results show that all columns 'Non-Null Count' are 10,000 which matches the number of rows in the shape shown above, meaning there are no missing values.

The 'Dtype' result also shows that all columns are either float(float 64) or integer(int64) datatype. The absence of the dtype 'object' signifies there are no categorical variables. See below

```
In [486]: dataset.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 50 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                ---
0   id                                    10000 non-null  int64
1   NumDots                              10000 non-null  int64
2   SubdomainLevel                      10000 non-null  int64
3   PathLevel                           10000 non-null  int64
4   UrlLength                           10000 non-null  int64
5   NumDash                             10000 non-null  int64
6   NumDashInHostname                   10000 non-null  int64
7   AtSymbol                            10000 non-null  int64
8   TildeSymbol                         10000 non-null  int64
9   NumUnderscore                       10000 non-null  int64
10  NumPercent                          10000 non-null  int64
11  NumQueryComponents                  10000 non-null  int64
12  NumAmpersand                       10000 non-null  int64
13  NumHash                             10000 non-null  int64
14  NumAtSymbol                         10000 non-null  int64
15  NumTildeSymbol                     10000 non-null  int64
16  NumUnderscoreInDomain               10000 non-null  int64
17  NumPercentInDomain                 10000 non-null  int64
18  NumQueryInDomain                   10000 non-null  int64
19  NumAmpersandInDomain                10000 non-null  int64
20  NumHashInDomain                     10000 non-null  int64
21  NumAtSymbolInDomain                 10000 non-null  int64
22  NumTildeSymbolInDomain              10000 non-null  int64
23  NumUnderscoreInPath                 10000 non-null  int64
24  NumPercentInPath                   10000 non-null  int64
25  NumQueryInPath                     10000 non-null  int64
26  NumAmpersandInPath                  10000 non-null  int64
27  NumHashInPath                       10000 non-null  int64
28  NumAtSymbolInPath                   10000 non-null  int64
29  NumTildeSymbolInPath                10000 non-null  int64
30  NumUnderscoreInPath                 10000 non-null  int64
31  NumPercentInPath                   10000 non-null  int64
32  NumQueryInPath                     10000 non-null  int64
33  AbnormalFormAction                  10000 non-null  int64
34  PctNullSelfRedirectHyperlinks       10000 non-null  float64
35  FrequentDomainNameMismatch          10000 non-null  int64
36  FakeLinkInStatusBar                 10000 non-null  int64
37  RightClickDisabled                  10000 non-null  int64
38  PopUpWindow                         10000 non-null  int64
39  SubmitInfoToEmail                   10000 non-null  int64
40  IframeOrFrame                       10000 non-null  int64
41  MissingTitle                        10000 non-null  int64
42  ImagesOnlyInForm                    10000 non-null  int64
43  SubdomainLevelRT                    10000 non-null  int64
44  UrlLengthRT                         10000 non-null  int64
45  PctExtResourceUrlsRT                10000 non-null  int64
46  AbnormalExtFormActionR              10000 non-null  int64
47  ExtMetaScriptLinkRT                10000 non-null  int64
48  PctExtNullSelfRedirectHyperlinksRT  10000 non-null  int64
49  CLASS_LABEL                         10000 non-null  int64
dtypes: float64(3), int64(47)
memory usage: 3.8 MB

In [487]: #From the result above, it can be seen that the dataset has no missing values since the non-null Count is equal to the number of
```

## Exploratory Data Analysis

In this stage, the '.describe' and '.value\_counts' function was employed to have a glimpse into the distribution of the X-columns (features to be used to define the class) values, and the Y-variable (the class). See below

```
In [488]: dataset.describe()

Out[488]:
```

	id	NumDots	SubdomainLevel	PathLevel	UrlLength	NumDash	NumDashInHostname	AtSymbol	TildeSymbol	NumUnderscc
count	10000.00000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	2.445100	0.586800	3.300300	70.264100	1.818000	0.138900	0.000300	0.013100	0.323
std	2886.89568	1.346836	0.751214	1.863241	33.369877	3.106258	0.545744	0.017319	0.113709	1.114
min	1.00000	1.000000	0.000000	0.000000	12.000000	0.000000	0.000000	0.000000	0.000000	0.000
25%	2500.75000	2.000000	0.000000	2.000000	48.000000	0.000000	0.000000	0.000000	0.000000	0.000
50%	5000.50000	2.000000	1.000000	3.000000	62.000000	0.000000	0.000000	0.000000	0.000000	0.000
75%	7500.25000	3.000000	1.000000	4.000000	84.000000	2.000000	0.000000	0.000000	0.000000	0.000
max	10000.00000	21.000000	14.000000	18.000000	253.000000	55.000000	9.000000	1.000000	1.000000	18.000

8 rows x 50 columns

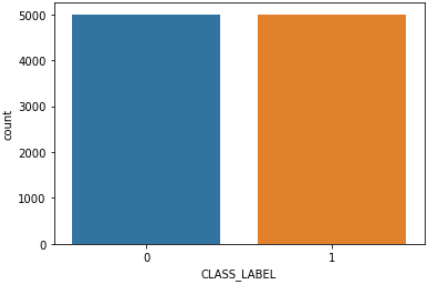
```
In [489]: #Doing some exploration on the Y-variable
dataset.CLASS_LABEL.value_counts()

Out[489]: 1    5000
          0    5000
          Name: CLASS_LABEL, dtype: int64
```

The graph below shows that the Y-variable is evenly distributed:

```
In [490]: #To view the Class_Label distribution
sns.countplot(data=dataset, x='CLASS_LABEL')

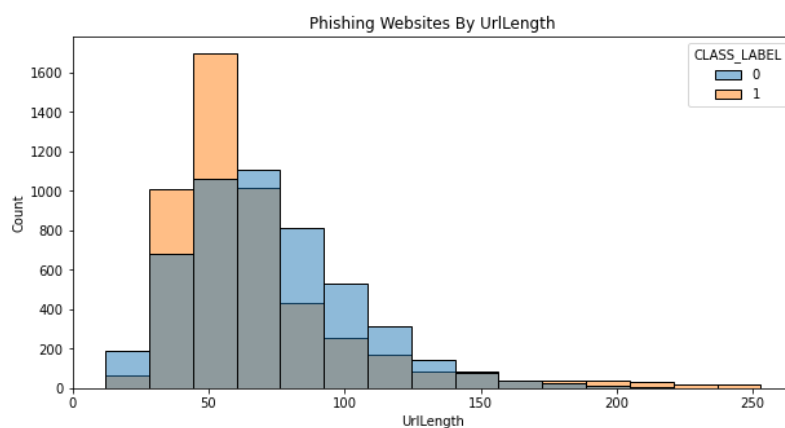
Out[490]: <AxesSubplot:xlabel='CLASS_LABEL', ylabel='count'>
```



```
In [491]: #From the above result, it can be immediately seen that the Class Label is split evenly between Legitimate and Phishing webpages
```

Next, some specific attributes were considered to check their distribution and proportion in the class label. Below:

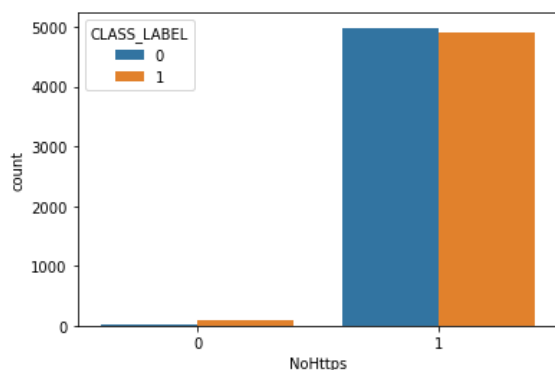
```
In [492]: plt.figure(figsize=(10,5))
plt.title('Phishing Websites By UrlLength')
sns.histplot(data=dataset, hue='CLASS_LABEL', x='UrlLength', bins=15)
plt.show()
```



From the distribution in the image above, pages with URL lengths of around 50 have the highest number of phishing links.

```
In [493]: #To see how the NoHttps attribute is distributed and its proportion in the CLASS_LABEL
sns.countplot(data=dataset, x='NoHttps', hue='CLASS_LABEL')
```

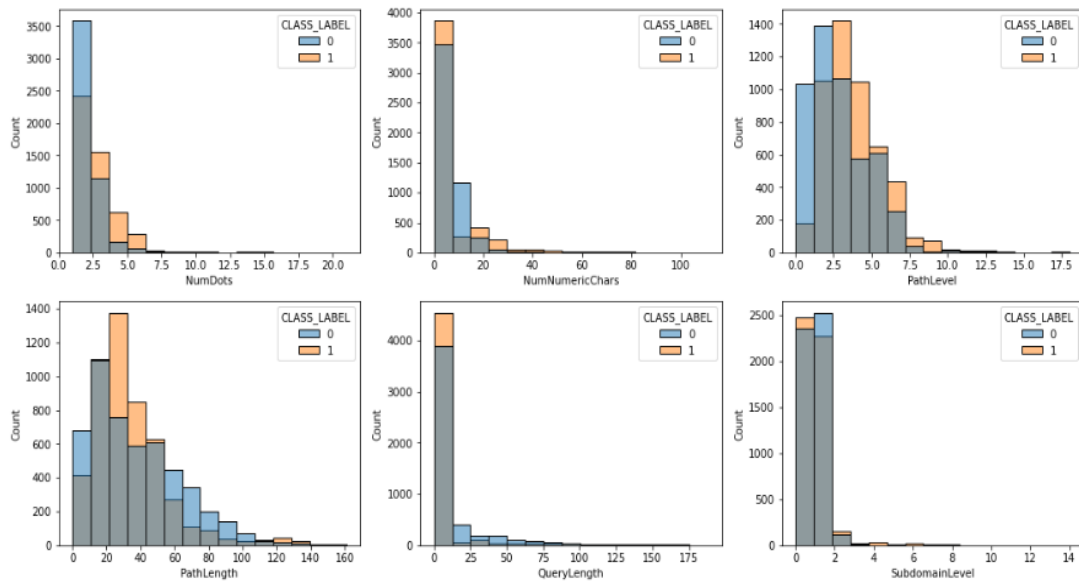
```
Out[493]: <AxesSubplot:xlabel='NoHttps', ylabel='count'>
```



Next, a function to show the histogram distribution between selected variables and the class label. See below:

```
#Defining a function to show me the histogram and boxplot views of some selected Variables
plot_Var = ['NumDots', 'NumNumericChars', 'PathLevel', 'PathLength', 'QueryLength', 'SubdomainLevel']
def chart_plot(plot_type):
    if plot_type=='histogram':
        plt.figure(figsize=(18,9))
        for i in range(len(plot_Var)):
            plt.subplot(2,3,i+1)
            sns.histplot(data=dataset, hue='CLASS_LABEL', x=plot_Var[i], bins=15)
    elif plot_type=='boxplot':
        plt.figure(figsize=(18,9))
        for i in range(len(plot_Var)):
            plt.subplot(2,3,i+1)
            sns.boxplot(data=dataset, hue='CLASS_LABEL', x=plot_Var[i], showmeans=True)
    else:
        pass
```

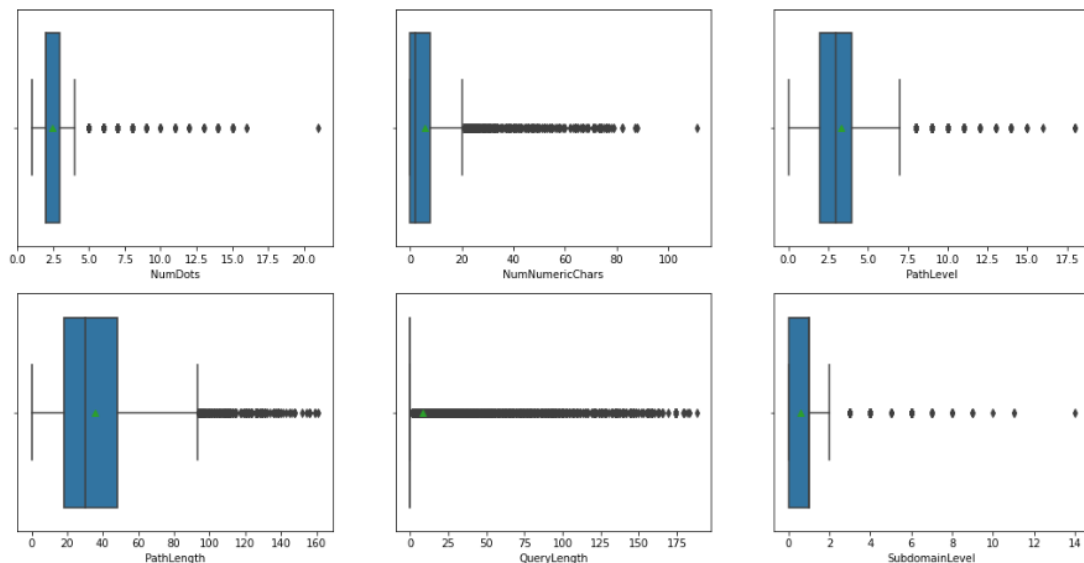
```
In [495]: #To quickly examine the distribution between some variables and the Class Label.
chart_plot('histogram')
```



The results above show that the distribution of most of the variables above is concentrated towards the lower values.

The boxplots below reveal the outliers in the above-selected variables.

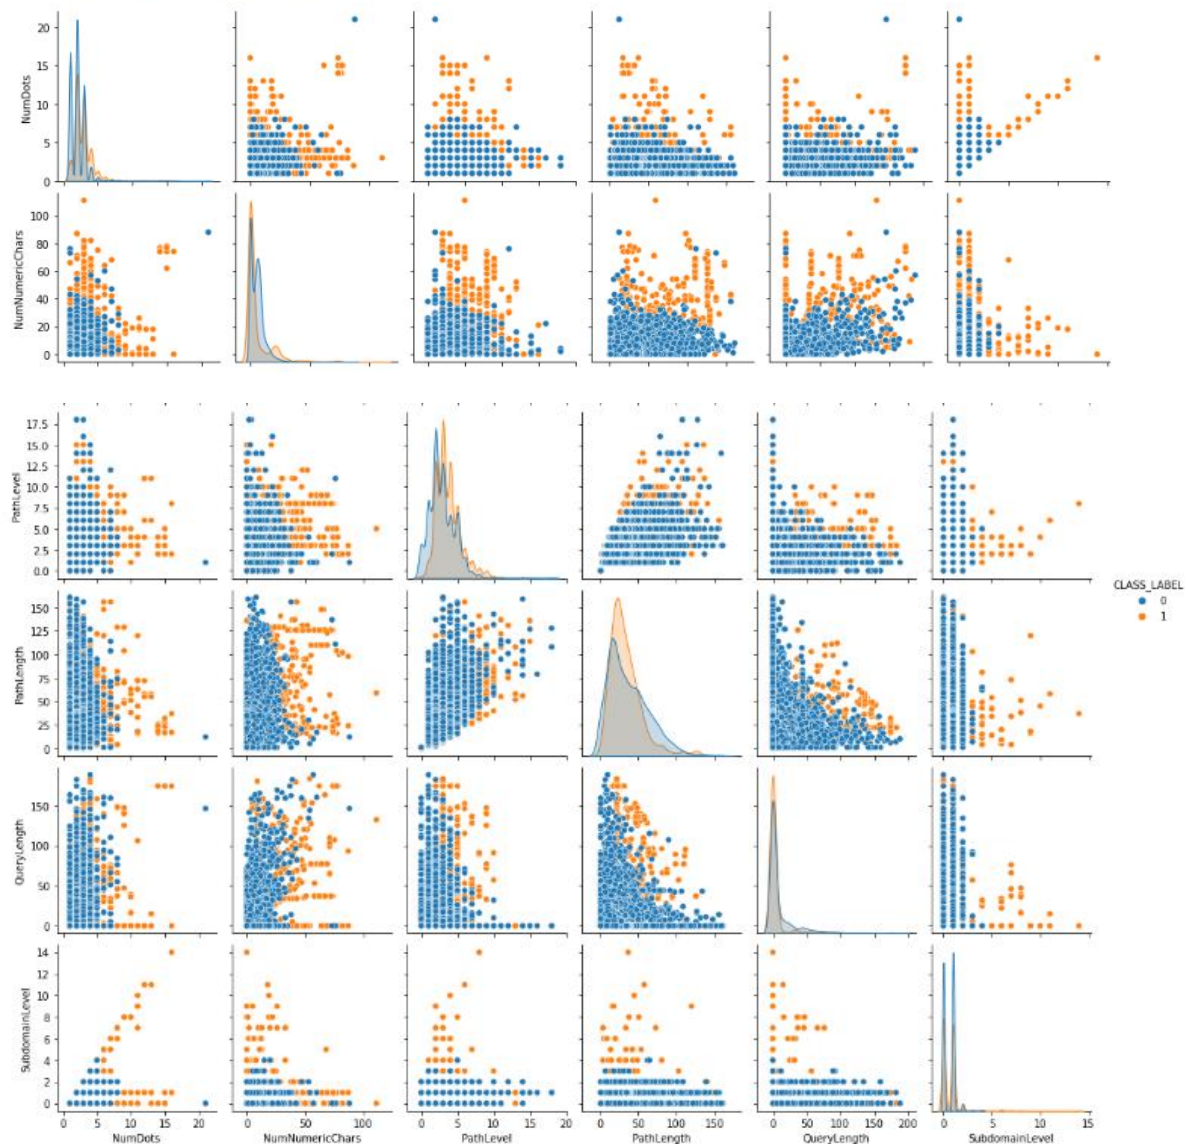
```
In [496]: #To have a quick view on outlier values using a boxplot
chart_plot('boxplot')
```



To get a view of how the variables are related to themselves and the class label, we plot a scatterplot matrix. The diagonal in the matrix shows the density plot of the label. View the code & result below:

```
In [497]: #To get the relationship between the above selected variables and the Class Label, we plot a scatterplot matrix
plt.figure(figsize=(15,10))
sns.pairplot(dataset[['NumDots', 'NumNumericChars', 'PathLevel', 'PathLength', 'QueryLength', 'SubdomainLevel', 'CLASS_LABEL']],
             hue='CLASS_LABEL')
plt.show()
```

<Figure size 1080x720 with 0 Axes>



## Feature Improvement

As shown earlier in **\*\*insert section here\*\***, this dataset does not contain categorical variables, therefore no encoding using LabelEncoder, One-hot Encoder, etc is necessary.

Also, there are no missing values. As a result, there was no need to employ SimpleImputer to impute missing values

The 'Id' column was dropped since it doesn't affect the classification and the class label was renamed to 'Result'.

## Feature Improvement

```
In [498]: dataset1 = dataset.drop('id', axis=1) #Drop the ID column (since it is not an attribute to be considered in the classification)
dataset1 = dataset1.rename(columns={'CLASS_LABEL':'Result'}) #Renaming the CLASS_LABEL column
dataset1.head()
```

```
Out[498]:
```

	NumDots	SubdomainLevel	PathLevel	UrlLength	NumDash	NumDashInHostname	AtSymbol	TildeSymbol	NumUnderscore	NumPercent	...	IframeOrFrame
0	3	1	5	72	0	0	0	0	0	0	...	0
1	3	1	3	144	0	0	0	0	2	0	...	0
2	3	1	2	58	0	0	0	0	0	0	...	0
3	3	1	6	79	1	0	0	0	0	0	...	0
4	3	0	4	46	0	0	0	0	0	0	...	1

5 rows x 49 columns



## Implementation in Python

### Brief description of the algorithms used

#### K-Nearest Neighbors (KNN) algorithm

The KNN algorithm is a python classification algorithm that works using a simple idea but provides great results. Usually, the algorithm uses the Euclidean distance,  $d$ , to find the distance between 2 points  $(x_1, y_1)$  and  $(x_2, y_2)$ .

$$d = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

To make a classification, the system uses observations from a training set, calculates the Euclidean distance between them, gets the  $k$  nearest neighbors (where  $k$  is the number of neighbors chosen), and uses the attributes of the neighbors to make a classification. (Baoli, Shiwen, & Qin, 2003)

#### Artificial Neural Networks

ANN classification is modeled loosely on the functionality of the human brain. ANN consists of processing nodes that are interconnected densely. These nodes are separated into layers called input, hidden, and output layers, and work by assigning weights to the input signals attributes, processing them, and then passing them on to other nodes in the next tier to eventually establish a classification. (Hardesty, 2017)

## Applications of the Algorithms

### K-Nearest Neighbors Classification

#### Application of the KNN algorithm to the dataset & Model design

First, the data columns were split into **X** and **y** representing the **Features** and the **Class** respectively.

Then, the datasets were split into training and testing data for both the **X** and the **y** parts using the `train_test_split` function from `sklearn.model_selection`.

The dataset was split into a 70:30 training-testing ratio, 'random\_state' integer was specified so that the same train and test sets are gotten across different executions.

Next, the `StandardScaler` function was used to fit and transform the **X** training and just transform the **X** testing dataset (to apply the initial fitting parameters to the test dataset).

Conducting the `fit_transform` statistically means centering the data by subtracting the mean and dividing by the standard deviation.

$$x' = (x - \mu) / \sigma$$

`n-jobs` was set to -1 to activate parallel processing for faster run-times

## k-Nearest Neighbors (KNN) Classification

```
In [75]: #Split the data columns into arrays containing attributes and another containing results
X = dataset1.iloc[:, 0: len(dataset1.columns.tolist())-1]
y = dataset1.iloc[:, -1]

In [76]: from sklearn.model_selection import train_test_split, GridSearchCV
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y) #Stratify keeps the split percentage constant in both training and testing

In [77]: #To fit transform the train values and transform the test values
sc=StandardScaler()
X_train_s=sc.fit_transform(X_train) #This returns a numpy array, and that is because numpy arrays are faster to process than dataframes
X_test_s=sc.transform(X_test)

In [78]: #Using the KNN classifier
from sklearn.neighbors import KNeighborsClassifier
class_model = KNeighborsClassifier(n_neighbors=3, metric='minkowski', p=2, n_jobs=-1) #n_jobs =-1 maximizes CPU capacity to make use of all cores
class_model.fit(X_train_s, y_train)

Out[78]: KNeighborsClassifier(n_jobs=-1, n_neighbors=3)
```

### Experimental procedure

For the model, as seen in the figure above, the initial parameters used for the first experiment are:

N\_neighbors = 3  
Metric = Minkowski  
P=2

The approach to validation used is the Hold-out method. This method was chosen because it is ideal for when there is a relatively large amount of data.

### Result Visualization:

The model was deployed on the test data. The code & results are as shown:

### Evaluating the Model Performance

```
In [79]: y_pred = class_model.predict(X_test_s)
print(y_pred)

[1 1 0 ... 0 1 1]


In [80]: acc=accuracy_score(y_test,y_pred)
print('accuracy: %.2f\n\n'%(acc))
cm=confusion_matrix(y_test,y_pred)
print('Confusion Matrix')
print(cm, '\n\n')
print('=====')
result=classification_report(y_test,y_pred)
print('Classification Report:')
print('-----')
print(result)
disp=ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_model.classes_)
disp.plot()
plt.show()
```

accuracy:0.96

Confusion Matrix  
[[1428 72]  
[ 48 1452]]

=====  
Classification Report:  
=====

	precision	recall	f1-score	support
0	0.97	0.95	0.96	1500
1	0.95	0.97	0.96	1500
accuracy			0.96	3000
macro avg	0.96	0.96	0.96	3000
weighted avg	0.96	0.96	0.96	3000



To set and optimize hyper-parameters, the use of GridSearchCV was employed, and a graph of the various k-values against their corresponding accuracy is shown below:

```
In [81]: knn2=KNeighborsClassifier()
hyperparameters={'n_neighbors': (1,12,1), 'metric': ('minkowski', 'chebyshev')}

knn_cv = GridSearchCV(knn2, hyperparameters, n_jobs=-1, verbose=1)
knn_cv.fit(X_train_s, y_train)
```

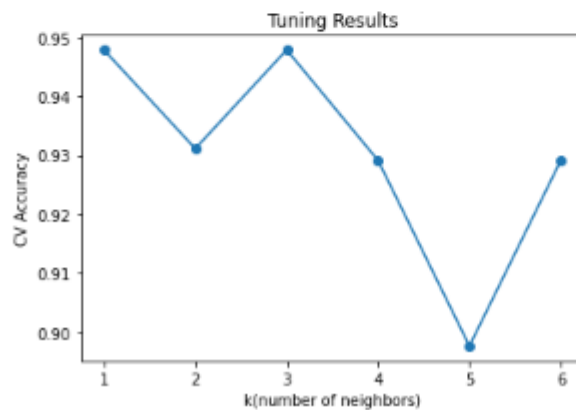
Fitting 5 folds for each of 6 candidates, totalling 30 fits

```
Out[81]: GridSearchCV(estimator=KNeighborsClassifier(), n_jobs=-1,
                      param_grid={'metric': ('minkowski', 'chebyshev'),
                                   'n_neighbors': (1, 12, 1)},
                      verbose=1)
```

```
In [82]: knn_cv.best_params_
```

```
Out[82]: {'metric': 'minkowski', 'n_neighbors': 1}
```

```
In [83]: tuning_result=knn_cv.cv_results_['mean_test_score']
plt.plot(range(1,7), tuning_result, 'o-')
plt.ylabel('CV Accuracy')
plt.xlabel('k(number of neighbors)')
plt.title('Tuning Results')
plt.show()
```



From the figure above, the best parameters are:

Metric: Minkowski

k: 1

The graph above and the tuning result below show that k=1 and k=3 yield the same accuracy score.

```
In [84]: tuning_result
```

```
Out[84]: array([0.94785714, 0.93114286, 0.94785714, 0.92914286, 0.89771429,
                0.92914286])
```


**Next experiment using derived parameters k=1**

Results Visualization below:

```
In [85]: predict_best = knn_cv.best_estimator_.predict(X_test_s)
print(classification_report(y_test, predict_best))
print("=====")
print("Estimated Value Confusion Matrix\n*****")
cm2 = confusion_matrix(y_test, predict_best)
disp=ConfusionMatrixDisplay(confusion_matrix=cm2, display_labels=class_model.classes_)
disp.plot()
plt.show()
```

	precision	recall	f1-score	support
0	0.96	0.95	0.96	1500
1	0.96	0.96	0.96	1500
accuracy			0.96	3000
macro avg	0.96	0.96	0.96	3000
weighted avg	0.96	0.96	0.96	3000

=====  
Estimated Value Confusion Matrix  
\*\*\*\*\*





Comparing both results (k=3, k=1) side by side:

=====  
Classification Report:

	precision	recall	f1-score	support
0	0.97	0.95	0.96	1500
1	0.95	0.97	0.96	1500
accuracy			0.96	3000
macro avg	0.96	0.96	0.96	3000
weighted avg	0.96	0.96	0.96	3000

	precision	recall	f1-score	support
0	0.96	0.95	0.96	1500
1	0.96	0.96	0.96	1500
accuracy			0.96	3000
macro avg	0.96	0.96	0.96	3000
weighted avg	0.96	0.96	0.96	3000

=====  
Estimated Value Confusion Matrix  
\*\*\*\*\*



From the classification reports, both results are evenly matched. K=1 is finally chosen as it has a more balanced overall score.

## Relevant Literature

In the journal article, "A Review of Data Classification Using K-Nearest Neighbour Algorithm" by Aman Kataria, M. D. Singh, the contributors alluded to the fact that even though the KNN model is efficient, it is limited by great calculation complexity, cost, operating speed, and complete dependence on the training set. The results of the experiment align with this allusion. (Kataria & Singh, 2013)

## Artificial Neural Networks (ANNs) classification

### Application of the ANN algorithm to the dataset

Since the process of splitting, fitting, and transforming the dataset is the same as the algorithm used above, the variables holding the features part of the training and testing data were copied.

```
In [86]: X_train_NN = X_train_s.copy()
        X_test_NN = X_test_s.copy()
```

### Explanation of the experimental procedure

First, the number of features (columns) was taken to estimate the number of features for the input layer:

```
In [87]: X_train_NN.shape[1] #To get the number of features for the input Layer
Out[87]: 48
```

Next, a set\_seed function was defined to make sure the results are reproducible each time the model is run. Random seeds were taken from Numpy, Python, and Tensorflow:

```
In [88]: #This step is to make sure the results can be reproduced and do not change each time the model is run
def set_seed(seed=20):
    np.random.seed(seed) #to seed from everything in numpy
    random.seed(seed) #to seed from everything in Python itself
    tf.random.set_seed(seed) #seeds randomness from anything in tensorflow
```

The Sequential API in Keras was used as the model is creating layers in sequence i.e input – Hidden – Output layers and no branching/sub-classes are required:

Next, the input shape of the hidden layer was specified, and other layers' shapes are generated automatically. 'Dense' is used because all neurons from a preceding layer feed into each neuron of the next layer. The activation function used is 'ReLU', and an arbitrary number of neurons (30) was selected.

In the Output layer, the Sigmoid activation function was employed as we only have 2 mutually exclusive classes.

```
In [89]: set_seed()
model = tf.keras.models.Sequential() #The Sequential here is the type of API we're using
model.add(tf.keras.layers.Dense(30,activation='relu',input_shape=(X_train_NN.shape[1],))) #Only the input shape needs to be spec
model.add(tf.keras.layers.Dense(1,activation='sigmoid')) #sigmoid is used for binary classifications where we only have 2 mutual
#when writing report, check how to determine the optimum number of input neurons ("20" in line 2)
model.summary()
#The number of parameters is gotten by multiplying the number of hidden layer nodes by the number of input features + bias (1 pe
```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 30)	1470
dense_3 (Dense)	(None, 1)	31

=====  
Total params: 1,501  
Trainable params: 1,501  
Non-trainable params: 0

### Model Parameters Used

Optimizer = 'adam'. This was used as it selects a good rate at which the weights are updated

Loss = 'binary\_crossentropy' This is because only 2 output classes are expected

Metrics = 'accuracy'. This was used because, as shown earlier, the output classes in the dataset are not imbalanced

Class weight wasn't specified as the output class is balanced.

```
In [90]: model.compile(optimizer='adam',  
                    loss='binary_crossentropy',  
                    metrics='accuracy') #binary_crossentropy Because there are only 2 classes, and accuracy metrics because my class i
```

```
In [91]: history = model.fit(X_train_NN, y_train,  
                            batch_size = 10,  
                            epochs= 30,  
                            verbose=2,  
                            validation_split=0.2) #didn't specify class weight because my class is balanced
```


Epoch 1/30  
560/560 - 4s - loss: 0.3010 - accuracy: 0.8875 - val\_loss: 0.1710 - val\_accuracy: 0.9400 - 4s/epoch - 7ms/step  
Epoch 2/30  
560/560 - 1s - loss: 0.1635 - accuracy: 0.9405 - val\_loss: 0.1328 - val\_accuracy: 0.9521 - 1s/epoch - 2ms/step  
Epoch 3/30  
560/560 - 1s - loss: 0.1381 - accuracy: 0.9471 - val\_loss: 0.1183 - val\_accuracy: 0.9600 - 985ms/epoch - 2ms/step  
Epoch 4/30  
560/560 - 1s - loss: 0.1240 - accuracy: 0.9511 - val\_loss: 0.1110 - val\_accuracy: 0.9593 - 960ms/epoch - 2ms/step  
Epoch 5/30  
560/560 - 1s - loss: 0.1106 - accuracy: 0.9571 - val\_loss: 0.0990 - val\_accuracy: 0.9614 - 965ms/epoch - 2ms/step  
Epoch 6/30  
560/560 - 1s - loss: 0.1009 - accuracy: 0.9611 - val\_loss: 0.0942 - val\_accuracy: 0.9636 - 969ms/epoch - 2ms/step  
Epoch 7/30  
560/560 - 1s - loss: 0.0915 - accuracy: 0.9659 - val\_loss: 0.0891 - val\_accuracy: 0.9664 - 982ms/epoch - 2ms/step  
Epoch 8/30  
560/560 - 1s - loss: 0.0847 - accuracy: 0.9684 - val\_loss: 0.0848 - val\_accuracy: 0.9693 - 982ms/epoch - 2ms/step  
Epoch 9/30  
560/560 - 1s - loss: 0.0781 - accuracy: 0.9723 - val\_loss: 0.0829 - val\_accuracy: 0.9679 - 954ms/epoch - 2ms/step  
Epoch 10/30  
560/560 - 1s - loss: 0.0735 - accuracy: 0.9730 - val\_loss: 0.0808 - val\_accuracy: 0.9707 - 949ms/epoch - 2ms/step  
Epoch 11/30  
560/560 - 1s - loss: 0.0691 - accuracy: 0.9750 - val\_loss: 0.0804 - val\_accuracy: 0.9679 - 1s/epoch - 2ms/step  
Epoch 12/30  
560/560 - 1s - loss: 0.0646 - accuracy: 0.9771 - val\_loss: 0.0799 - val\_accuracy: 0.9721 - 1s/epoch - 2ms/step  
Epoch 13/30  
560/560 - 1s - loss: 0.0620 - accuracy: 0.9764 - val\_loss: 0.0788 - val\_accuracy: 0.9721 - 1s/epoch - 2ms/step  
Epoch 14/30  
560/560 - 1s - loss: 0.0576 - accuracy: 0.9798 - val\_loss: 0.0794 - val accuracy: 0.9707 - 1s/epoch - 2ms/step

## Evaluating the Neural Network

The training and validation accuracy shows an initial steep increase followed by a gradual increase in accuracy for the training set and a fluctuating constant for the validation set. See below:

```
In [92]: accuracy = history.history['accuracy']  
validation_accuracy = history.history['val_accuracy']  
plt.plot(accuracy, label='Training Set Accuracy')  
plt.plot(validation_accuracy, label='Validation Set Accuracy')  
plt.ylabel('Accuracy')  
plt.ylim([min(plt.ylim()),1])  
plt.title('Training and Validation Accuracy Across Epochs')  
plt.legend()
```


Out[92]: <matplotlib.legend.Legend at 0x25d866c0ee0>



The figure below shows an initial steep decline followed by a gradual decline in both the training and accuracy set loss

```
In [93]: loss = history.history['loss']
validation_loss = history.history['val_loss']
plt.plot(loss, label='Training Set Loss')
plt.plot(validation_loss, label='Validation Set Loss')
plt.ylabel('Loss')
plt.title('Training and Validation Accuracy Across Epochs')
plt.legend()
```

```
Out[93]: <matplotlib.legend.Legend at 0x25d86712190>
```



The model was used on the test data set and the results were rounded. If 0.5 and above, it returns 1, if less than 0.5, it returns 0. Then the classification results report was printed. See images below

```
In [94]: model.evaluate(X_test_NN, y_test)
```

```
94/94 [=====] - 0s 2ms/step - loss: 0.0788 - accuracy: 0.9737
```

```
Out[94]: [0.07876517623662949, 0.9736666679382324]
```

```
In [95]: y_pred = model.predict(X_test_NN)
```

```
94/94 [=====] - 0s 2ms/step
```

```
In [96]: nn_pred = np.round(model.predict(X_test_NN)) #to convert the probability results to binary output where less than 0.5 is 0, greater than 0.5 is 1
```

```
94/94 [=====] - 0s 2ms/step
```

```
In [97]: print(classification_report(y_test, nn_pred))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	1500
1	0.97	0.97	0.97	1500
accuracy			0.97	3000
macro avg	0.97	0.97	0.97	3000
weighted avg	0.97	0.97	0.97	3000

## Relevant Literature

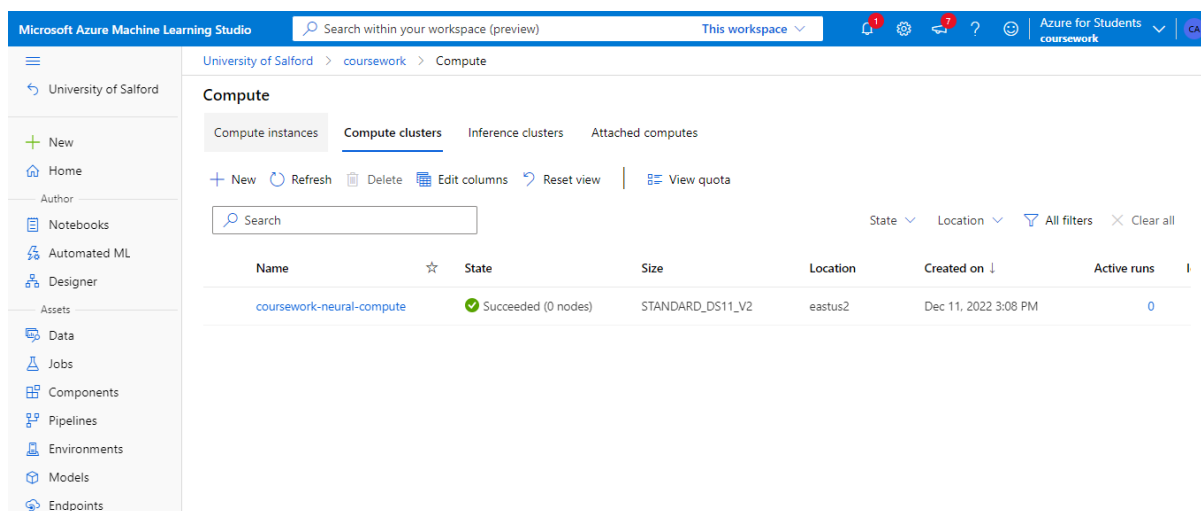
In the journal, “Basic Tenets of Classification Algorithms K-Nearest-Neighbor, Support Vector Machine, Random Forest and Neural Network: A Review”, the authors posit that even though Neural Networks are generally more accurate than kNN models, the latter is more commonly used as it is easier to configure and implement. (Boateng, Otoo, & Abaye, 2020)



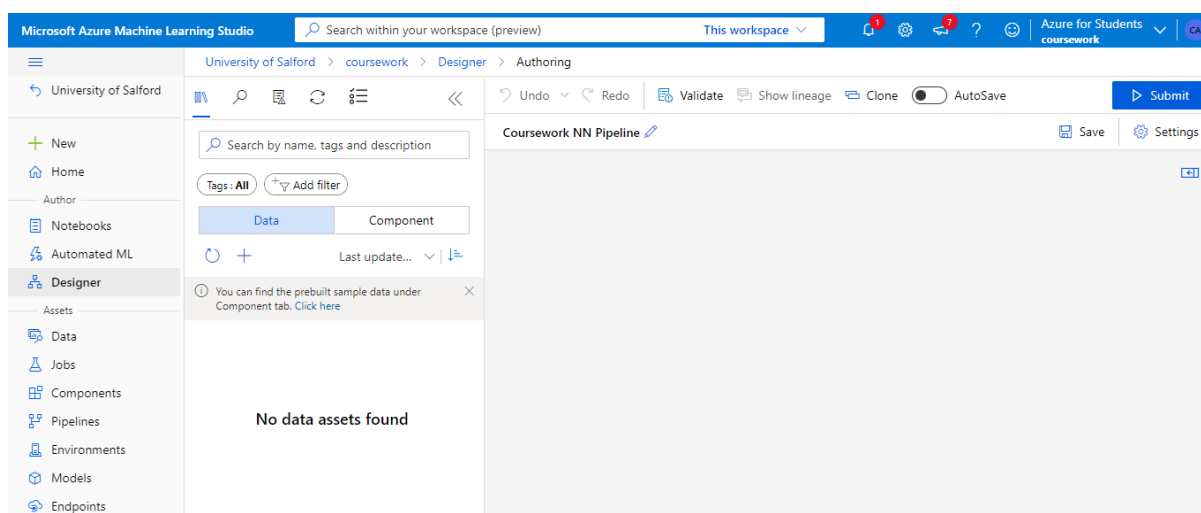
This experiment's results support the author above in showing that the Neural Network has more accuracy than the k-NN model.

## Implementation in Azure Machine Learning Studio

First, the compute cluster was created upon which the implementations would be run.



Next, a pipeline was created for the classification task. See below:



Afterward, a data asset was created by first specifying the Name, description, and type as a table, then the data source was specified as a local file, the storage type was in Azure blob storage, then the data was uploaded, and then a preview of the dataset is shown:

Microsoft Azure Machine Learning Studio

Search within your workspace (preview) This workspace

Azure for Students coursework

### Create data asset

- Data type
- Data source
- Storage type
- File or folder selection
- Settings**
- Schema
- Review

**Settings**  
These settings determine how the data is parsed. The initial settings are automatically detected; you can change them as needed to reparse the data.

File format: Delimited  
Delimiter: Comma  
Example: Field1,Field2,Field3  
Encoding: UTF-8

Column headers: All files have same headers  
Skip rows: None

☐ Dataset contains multi-line data

**Data preview**  
Processing tabular files with multi-line data is slower because multiple CPU cores cannot be used to ingest the data in parallel. Checking this option may result in slower

id	NumDots	Subdo...	PathLe...	UrlLen...	NumDa...	NumDa...	AtSym...	TildeSy...	NumU...	NumPe...	NumQ...
1	3	1	5	72	0	0	0	0	0	0	0
2	3	1	3	144	0	0	0	0	2	0	2
3	3	1	2	58	0	0	0	0	0	0	0
4	3	1	6	79	1	0	0	0	0	0	0

Back Next Review Cancel

A review of the data was checked before the data was created, and the result is shown below:

Review the settings for your data asset and make any changes as needed.

**Data type**

**Name**  
phishing\_dataset

**Description**  
This is a data containing information about web URLs that can be used to check if its a legitimate or a suspicious website

**Type**  
tabular

**Data source**

**Type**  
Local

**Schema**

id	Integer
NumDots	Integer
SubdomainLevel	Integer
PathLevel	Integer
UrlLength	Integer

(showing 5 of 51 columns)

**File selection**

**Upload path**  
azureml://subscriptions/8c5e200f-d252-4eeb-b4f1-9d3186ac5674/resourcegroups/Chike\_Azure\_Resource/workspaces/coursework/datastores/workspaceblobstore/paths/UI/2022-12-11\_161503\_UTC/Phishing\_Legitimate\_full.csv

**Files uploaded**  
Phishing\_Legitimate\_full.csv

**Storage**

**Datastore type**  
AzureBlob

**Datastore name**  
workspaceblobstore

**Settings**

**Delimiter**  
Comma

A quick look at the schema also gives some valuable information about the dataset like the datatypes for each column. It was also carefully scrutinized to ensure all datatypes were correct. The top part is shown below:

Microsoft Azure Machine Learning Studio

Search within your workspace (preview) This workspace

Azure for Students coursework

### Create data asset

- Data type
- Data source
- Storage type
- File or folder selection
- Settings
- Schema
- Review

Include	Column name	Type	Example values	Date format	Properties
<input type="checkbox"/>	Path	String		Not applicable to s...	Not applicable t...
<input checked="" type="checkbox"/>	id	Integer	1, 2, 3	Not applicable to s...	Not applicable t...
<input checked="" type="checkbox"/>	NumDots	Integer	3, 3, 3	Not applicable to s...	Not applicable t...
<input checked="" type="checkbox"/>	SubdomainLevel	Integer	1, 1, 1	Not applicable to s...	Not applicable t...
<input checked="" type="checkbox"/>	PathLevel	Integer	5, 3, 2	Not applicable to s...	Not applicable t...
<input checked="" type="checkbox"/>	UrlLength	Integer	72, 144, 58	Not applicable to s...	Not applicable t...
<input checked="" type="checkbox"/>	NumDash	Integer	0, 0, 0	Not applicable to s...	Not applicable t...
<input checked="" type="checkbox"/>	NumDashInHostname	Integer	0, 0, 0	Not applicable to s...	Not applicable t...

Back Next Cancel

The final data asset details are shown below:

Microsoft Azure Machine Learning Studio

Search within your workspace (preview) This workspace

Azure for Students coursework

University of Salford > coursework > Data > phishing\_dataset

phishing\_dataset Version: 1 (latest) ☆

Details Consume Explore Models Jobs

New version Refresh Generate profile Archive

**Attributes**

Type (mltable)  
Table (mltable)

Dataset type (from Azure ML v1 APIs)  
Tabular

Created by  
Chike Ayogu

Profile  
[View profile](#)  
Job: --

Files in dataset  
1

Total size of files in dataset  
1.341 MiB

Current version

**Tags**

No data

**Description**

This is a data containing information about web URLs that can be used to check if its a legitimate or a suspicious website

**Data sources**

Datatore: workspaceblobstore  
[UI/2022-12-11\\_161503\\_UTC/Phishing\\_Legitimate\\_full.csv](#)

## Pipeline Population

The imported dataset was added to the pipeline canvas and was previewed

Microsoft Azure Machine Learning Studio

Search within your workspace (preview) This workspace

Azure for Students coursework

University of Salford > coursework > Designer > Authoring

phishing\_dataset Version: 1 (latest) ☆

Details Consume Explore Models Jobs

New version Refresh Generate profile Archive

Attributes

Type (mltable)  
Table (mltable)

Dataset type (from Azure ML v1 APIs)  
Tabular

Created by  
Chike Ayogu

Profile  
[View profile](#)  
Job: --

Files in dataset  
1

Total size of files in dataset  
1.341 MiB

Current version

Tags

No data

Description

This is a data containing information about web URLs that can be used to check if its a legitimate or a suspicious website

Data sources

Datatore: workspaceblobstore  
[UI/2022-12-11\\_161503\\_UTC/Phishing\\_Legitimate\\_full.csv](#)

Coursework NN Pipeline

phishing\_dataset phishing\_dataset

DataOutput

Preview Profile

Number of columns: 50 Number of rows: First 50

id	NumDots	Subdo...	PathLe...	UrlL
1	3	1	5	72
2	3	1	3	144
3	3	1	2	58
4	3	1	6	79
5	3	0	4	46
6	3	1	1	42
7	2	0	5	60

Since there was no missing data in the initial dataset, there was no need to use the 'clean missing data' module.

Next, the 'split Data' module was used to split the data into 70% training and 30% testing. Randomized split was set to True, and the random seed 42, the stratified split was set to true to keep the percentages constant, and the stratified column was set to my target variable.

### Split Data

Splitting mode ⓘ \*

Split Rows

Fraction of rows in the first output dataset ⓘ \* ...

0.7

Randomized split ⓘ \*

True

Random seed ⓘ \* ...

42

Stratified split ⓘ \*

True

Stratification key column ⓘ \* [Edit column](#)

Column names: CLASS\_LABEL

Next, the train model module was selected and the label column was specified as 'CLASS\_LABEL'

### Train Model

Label column ⓘ \* [Edit column](#)

Column names: CLASS\_LABEL

Model explanations ⓘ

False

---

Output settings >

---

Run settings >

---

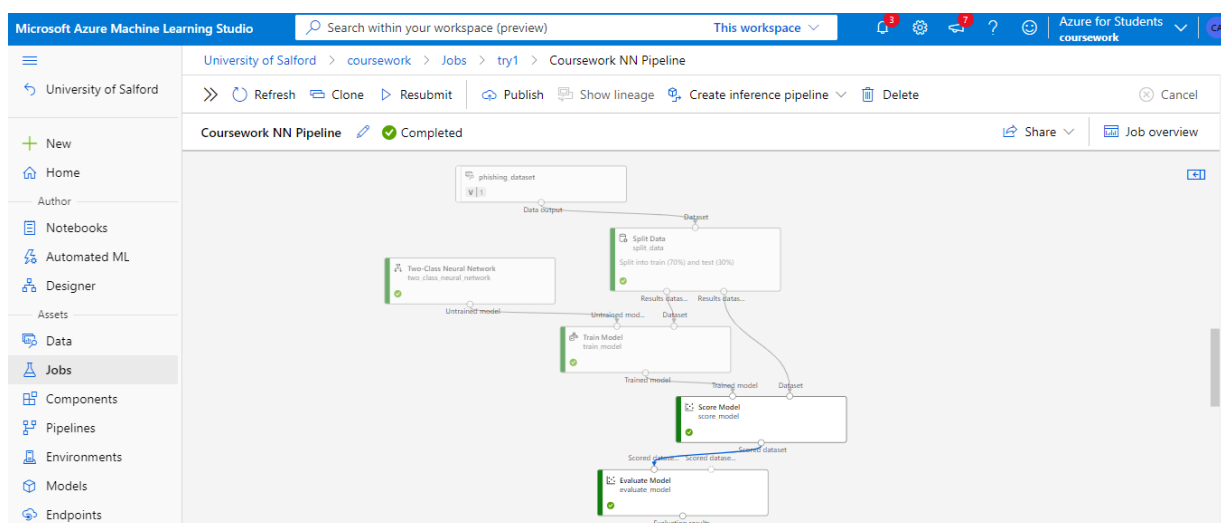
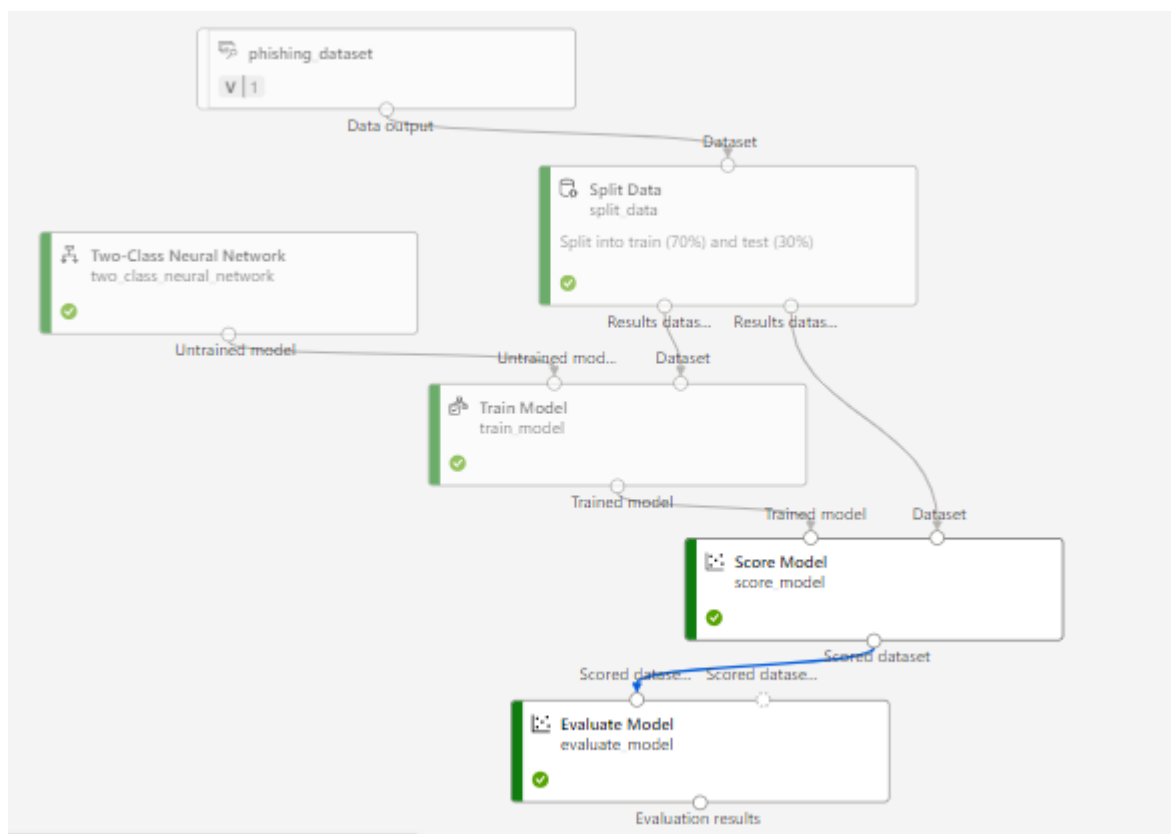
Node information >

---

Component information >

Afterward, a 'Two-class Neural Network' module was introduced, the number of hidden nodes was specified to 50, and was connected to the pipeline. Then, the 'Evaluate model' module was added, and the entire pipeline run.

The final pipeline is shown below:




When the pipeline was run, a new column was appended at the end with the predicted values. A sample is shown below:

Rows ⑦  
3,000

Columns ⑦  
52

tHyper	CLASS_LABEL	Scored Labels	Scored Probabilities
	0	0	0.000045
	1	1	0.999604
	1	1	0.999274
	1	1	0.999814
	1	1	0.999999
	0	0	0.000301
	1	1	0.99035
	0	0	0.000201
	0	0	0.000001
	1	1	0.998609
	0	0	0.00015



To view, select a column in the table

Close examination of the results of the 'evaluation model' module shows the results score and the classification model as seen below.

Threshold  0.5

Accuracy 0.995  
Precision 0.996  
Recall 0.993  
F1 Score 0.995  
AUC 1

		Actual	
		1	0
Predicted	1	1 490	6
	0	10	1 494

## Results analysis and discussion

### Performance metric used

The main performance metric used in evaluating both models is Accuracy. This was used because the dataset class output is well-balanced at 50:50.

A Confusion Matrix was also employed to have a broader look at the results.

### Presentation of results

#### KNN Classification

The confusion matrix and classification report using:


Randomly chosen (k=3). Below:

```
accuracy:0.96
```

```
Confusion Matrix  
[[1428  72]  
 [ 48 1452]]
```

```
=====  
Classification Report:  
-----
```


	precision	recall	f1-score	support
0	0.97	0.95	0.96	1500
1	0.95	0.97	0.96	1500
accuracy			0.96	3000
macro avg	0.96	0.96	0.96	3000
weighted avg	0.96	0.96	0.96	3000



Derived best parameter (k=1) is shown below:


	precision	recall	f1-score	support
0	0.96	0.95	0.96	1500
1	0.96	0.96	0.96	1500
accuracy			0.96	3000
macro avg	0.96	0.96	0.96	3000
weighted avg	0.96	0.96	0.96	3000

\*\*\*\*\*  
 Estimated Value Confusion Matrix  
 \*\*\*\*\*



The figure below shows the graph of the best parameters and their corresponding accuracy after hyperparameter tuning



k(Number of Neighbors)	Accuracy
1	0.94785714
2	0.93114286
3	0.94785714
4	0.92914286
5	0.89771429
6	0.92914286



### Artificial Neural Network

View graphs of Training and Validation set accuracy and losses below:





Classification report of the neural network:

```
In [97]: print(classification_report(y_test, nn_pred))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	1500
1	0.97	0.97	0.97	1500
accuracy			0.97	3000
macro avg	0.97	0.97	0.97	3000
weighted avg	0.97	0.97	0.97	3000

## **Azure Machine Learning**

The scores and the classification reports are seen below:

Threshold  0.5

Accuracy 0.995  
Precision 0.996  
Recall 0.993  
F1 Score 0.995  
AUC 1

		Actual	
		1	0
Predicted	1	1 490	6
	0	10	1 494

### Comparison and discussion of results

A side-by-side comparison of the results of the K-Nearest Neighbors and Neural Networks algorithms has been drafted, and is shown below:

	k-Nearest Neighbors	Neural Network
Accuracy	0.959333	0.973667
Precision	0.955086	0.973351
Recall	0.964000	0.974000
f1	0.959522	0.973675

From the results shown above, the Neural Network performs better than the K-NN algorithm judging by the listed performance metrics. Its accuracy, which was the main focus, was greater than that of the kNN by 1.43%. It also surpassed the k-NN model in Precision, Recall, and F1-Score by 1.83%, 1%, and 1.42% respectively.

The Azure machine learning model gives the best results among all three with an accuracy score of 99.5%

This metric (accuracy) measures the percentage of correctly predicted values (both phishing and legitimate websites) to the true values.

k(Number of Neighbors)	Accuracy
1	0.94785714
2	0.93114286
3	0.94785714
4	0.92914286
5	0.89771429
6	0.92914286

It can also be seen from the 'k-Neighbors against accuracy' table above, a slight change in the k-number can lead to a drastic change in the accuracy of the model. This implies that the model might not be suitable for future use with an entirely different dataset.

## **Ethical, legal, and professional considerations**

Some considerations made in the

1. The dataset poses no risks to individuals or organisations by making sure the URL names were excluded.
2. The data involved was collected from legitimate, publicly available means.
3. All parties involved in the collection of the dataset were duly listed

## **Conclusion**

Phishing attacks are a plague in this computer-driven world and these classification and prediction models with high accuracy can curb this scourge. All models showed good performance with the Neural Network having a slight edge over the k-NN model, and the Azure machine learning model performing best. This does not guarantee that these lofty scores will be replicated when applied to future datasets, however, it shows how useful these classification models can be in solving the phishing menace.